

1. Traceability

- Every assistant action, claim, or component field is traceable to a user message, prior context, or a tool output
- No hallucinated data, files, entities, or side effects
- Every action that requires a tool has a corresponding tool call

2. Tool Correctness

- All tools used are clearly defined in the sample
- Tool definitions are consistent and used uniformly throughout
- Tools are not used for tasks the model can do natively (e.g. formatting, parsing)
- Tool execution flags (e.g. `tool_executed = true`) are accurate

3. Temporal / Turn Logic

- Conversation flow is valid and logical
- When tools are used, the flow is: UP → Assistant → Tool Call → Tool Output → Assistant
- Sequential tool calls appear in separate assistant messages
- No user message directly follows a tool output

4. Component Sanity

- Components contain no clickable-looking elements or action buttons
- Components are generalizable and not overfit to a single scenario
- Component content is grounded in prior conversation context
- No unnecessary nesting or overly bespoke schemas

5. Evaluation Alignment

- Grading guidance is adapted per user turn
- Expected components are described in natural language (not strict lists)
- Expected components are empty when none should appear
- Clarifying-question turns are explicitly reflected in grading guidance